

Method: Extracting Calibrated Probability from Embeddings

Bare claim

"The experiment will succeed"

Embed

$\mathbf{e}_{bare}$

$$\Delta = \mathbf{e}_{hedged} - \mathbf{e}_{bare}$$

$$p = \Delta \cdot \mathbf{w}_{axis}$$

70.2%

Project onto
probability axis

Calibrated
probability

Hedged sentence

"The experiment will probably succeed"

Embed

$\mathbf{e}_{hedged}$

Difference vector